

Legally Subjective: How Much U.S. Supreme Court Vote Predictability Survives a Strong Ideological Baseline?

A Pre-Registered, Zero-Budget Experiment on Public Data

Amine Harch el Korane¹

¹Independent researcher. Project and code: github.com/Vitalcheffe/legally-subjective

Working paper, revision 2 · Corpus-Monde v1 (frozen 2026-08-28)

Abstract. Can a language model predict how a U.S. Supreme Court justice will vote from the public case file alone, without the verdict? And if the model is taught one justice’s own writing — a persona distilled from their past opinions — does it improve on that justice’s *future* cases? We frame both outcomes as publishable results and report the first two milestones of a pre-registered experiment conducted entirely on public data at zero cost. We assemble a frozen, hash-chained corpus of 569 argued cases (OT2015–OT2023, 13 justices, 4,569 coded justice×case votes) by fusing CourtListener metadata with Supreme Court Database votes. Before any model is trained, statistical baselines fix the price of admission: predicting each justice’s vote as their modal ideological direction from the training window achieves **63.7%** accuracy (Wilson 95% CI [61.3; 65.9]) at the vote level on the held-out window — the bar. We then train structured-feature challengers (logistic regression, gradient boosting, and interaction models over issue area, lower-court disposition, term, and argument length, with 50 sealed 5–4 cases excluded from every split): none beats the bar (58.6–60.4%; McNemar $p \leq 0.047$). A feature ablation shows justice identity is the only stable signal (−8.6 points without it). We are explicit about what this null does and does not establish: the *tested* structured representation does not beat a per-justice ideological baseline under this temporal protocol — it does not show that all non-textual information is exhausted, nor that the residual signal, if any, is textual. The remaining pre-registered conditions (zero-shot, persona fine-tuning, retrieval) probe the written record against a sealed holdout of fifty 5–4 decisions, where the ideological baseline is re-fit and re-evaluated on that population itself as the like-for-like comparator.

Keywords: judicial prediction; Supreme Court; pre-registration; supervised learning; attitudinal model; open data.

1 Introduction

Supreme Court votes are the most consequential binary predictions in American public life, and they are made under conditions that should humiliate any forecaster: the inputs are thousands of pages of legal argument, the decision rule is nine human beings, and the outcome distribution is close enough to a coin flip that luck is difficult to rule out. Yet a long line of work has shown that the task is not entirely random. Statistical models built on case attributes have predicted affirm/reverse outcomes at rates far above chance for two decades [1, 2], and the dominant theoretical account of the Court—the attitudinal model [3] — holds that ideology, not legal doctrine, explains most of the variance in votes.

This paper reports the design and the first results of *Legally Subjective*, an independently conducted, zero-budget experiment that asks a narrower and, we argue, more honest version of the prediction question. Rather than asking how accurately a court’s output can be retrofitted by a flexible learner with centuries of data, we ask two questions that can both fail loudly. First, is each justice’s vote predictable

from the public case file alone, on a fixed, modern window, against a pre-registered bar? Second—and this is the experiment’s live wire—if a language model is fine-tuned on one justice’s own past opinions, does it predict that justice’s *future* votes better than the same model without the persona? Both answers are publishable. If the persona helps, a *justice-specific textual signal* exists: the written voice carries predictive information about future decisions that ideology alone does not encode. Whether that signal is personality, drafting style, subject-matter specialization, or something else entirely is a distinct question this design does not settle—persona improvement is evidence of a signal, not of its nature. If it does not, then whatever is predictable about a justice on this window is already captured by their ideological direction, and the remainder is either private deliberation or noise.

Three design commitments distinguish the project from prior work. First, everything is pre-registered before the models see the *sealed* holdout: the baselines, the split, the metrics, and the fifty-case test set are frozen in the repository, with the selection seed and the SHA-256 of the sealed list published in advance. There is no second holdout and no post hoc tuning. Second, the experiment is built to be audited by a stranger: the corpus is assembled from public sources with every raw file hash-chained, and every number in this paper regenerates from scripts in the repository. Third, the budget is zero by construction—public data, free API tiers, and free-tier compute. We present this strictly as a reproducibility and accessibility constraint, not as evidence of scientific quality: it bounds what can be tested (model scale, hyper-parameter search, number of seeds) and it makes every result replicable by any researcher with a laptop, but a larger budget would not weaken the design—only widen the set of testable conditions.

The contributions of this working paper are: (i) a frozen, provenance-complete corpus of 569 argued Supreme Court cases spanning OT2015 through OT2023, fusing CourtListener metadata with Supreme Court Database (SCDB) vote coding, with 98.6% oral-argument coverage for future multimodal conditions; (ii) pre-registered statistical baselines, of which the strongest—per-justice modal ideological direction—sets a vote-level bar of 63.7% on the held-out window; (iii) a trained challenger, a family of structured-feature models whose *null result* is itself informative: nothing in the pre-decision metadata adds stable signal beyond who the justice is (Section 5); and (iv) the sealed, hash-frozen test protocol under which the remaining conditions (zero-shot, persona fine-tuning, retrieval) will eventually be judged (Section 4).

We emphasize what this paper does not claim. It does not claim that Supreme Court votes are unpredictable; two decades of literature and our own baselines say otherwise. It does not report any language-model results; the opinion-text corpus required for those conditions is still being collected under free-tier API quotas. And it does not treat the 63.7% bar as trivial: as Section 5 shows, three families of supervised models failed to clear it on the same rows, which is precisely why the bar is interesting.

2 Related work

Statistical forecasting of the Court. The modern line begins with the Supreme Court Forecasting Project: Ruger et al. [1] built a conditional-choice model with input from political science theory and, over a single term, predicted 75% of the Court’s affirm/reverse outcomes—against 59.1% for a panel of legal experts. The lesson was bruising for doctrine: a small regression with ideological covariates beat trained lawyers reading the briefs. A decade later, Katz et al. [2] scaled the exercise with a random forest over 628 features and nearly two centuries of decisions (1816–2013), reporting 70.2% accuracy at the case level and 71.9% at the individual vote level, and concluding that the Court behaves in a way that is, in their words, remarkably consistent with a machine-learning account. Both efforts, however, evaluate *in sample across history*: the models see several centuries — including eras with much higher unanimity rates—and predict votes whose contemporaneous context is far removed from today’s Court.

The attitudinal model. Political science supplies the theoretical prior for why such models work. Segal and Spaeth [3] argue that justices decide largely as attitudes wrapped in legal reasoning: revealed ideology dominates doctrine, and case facts matter mostly as triggers for those attitudes. Methodologically, Martin and Quinn [4] estimate dynamic ideal points for every justice since 1953, giving the field a continuous measure of the very covariate that our strongest baseline — per-justice modal direction — approximates with a single bit per justice. Our results in Section 5 are, in effect, a modern, pre-registered replication of the attitudinal prior on a nine-term window: the ideology bit is the signal; the structured case context is not.

Authorship and style. The persona condition is a stylometry question in modern dress. Mosteller and Wallace [5] settled the authorship of the disputed *Federalist* papers with function-word frequencies; the field they founded has since shown that individual writing style is persistent, measurable, and hard to disguise [6]. If a justice’s voice is a stable individual signature, fine-tuning a language model on their opinions should transfer to their future opinions; the open question—the one this experiment isolates — is whether that stylistic footprint carries *decision-relevant* information beyond the ideology bit, or merely rhetorical habits.

Language models and efficient fine-tuning. The tooling for the persona condition is recent. QLoRA [7] makes full-model fine-tuning tractable on free single-GPU tiers by quantizing the base model and training low-rank adapters, and open-weight families such as Llama 3 [8] provide a competent base. This is what converts the persona question from an institutional project into a zero-budget one: nine adapters, one per justice, trained on free-tier compute, each sealed alongside the data that produced it.

Positioning. Against this background, the present project narrows the claim in three ways that we believe matter for honesty of the result. The window is short and modern (OT2015–OT2023, an era of narrow majorities), so historical unanimity cannot inflate accuracy. The evaluation is pre-registered and forward-holding: a sealed set of fifty 5–4 decisions, chosen deterministically before any model touched them, is the only arbiter of the final conditions. And the bar is fixed *before* the challengers run, at the vote level, where the attitudinal bit is strongest—63.7% on our window — so that any reported improvement is an improvement over ideology itself, not over a straw man.

3 Data: Corpus-Monde v1

3.1 Sources and fusion

The corpus fuses two public sources, each chosen for a different epistemic role. *CourtListener* (Free Law Project) [9] supplies the case universe: dockets, dates, parties, opinion clusters, and oral-argument audio with transcripts, harvested from the 2026-06-30 bulk release and cross-checked against the REST API. The *Supreme Court Database* (SCDB), 2025_01 release, justice-centered case citation file [10], supplies the ground truth: per-justice vote directions coded by a single, long-running editorial team whose coding conventions are public and whose known limitations are documented in the database itself. CourtListener defines *what happened procedurally*; SCDB defines *what we count as a liberal or conservative vote*.

The corpus rule is deliberately narrow: every case argued before the Court in the window from 1 October 2015 through 30 June 2024 (terms OT2015 through OT2023) is included if it was argued and decided with a full opinion. Per curiams, memorandum decisions, and cases decided without argument are excluded—not because they are uninformative, but because the prediction task is defined over argued cases whose written record a model could read. Applying the rule yields 569 cases. Fusion keys on docket number with sub-docket-aware matching (a nontrivial step: consolidated dockets appear in SCDB as comma-separated lists and occasionally in variant forms such as original-jurisdiction numbering), and every matched pair is checked for term and case-name agreement.

Table 1: Corpus-Monde v1 at a glance (frozen 2026-08-28).

Quantity	Value
Argued cases, OT2015–OT2023	569
Justices appearing on the bench	13
Opinions (majority, dissent, concurrence)	1,778
Justice×case votes, coded direction	4,569
Cases with SCDB vote coding	96.1%
Cases with oral-argument audio + transcript	98.6%
5–4 decisions	79
of which sealed for the Final Test	50
Window	2015-10-01 to 2024-06-30




Figure 1: The corpus window. Each tick is one argued case, grouped by term (counts below each group as argued · 5–4). Red ticks are the 79 decisions that came down 5–4; the experiment’s window is deliberately the modern era of narrow majorities.

3.2 Freeze and provenance

The corpus is frozen: the identity of the data is fixed by a rule, and every artifact carries the hashes of everything upstream. The SHA-256 of every raw source file (bulk archives, API pages, the SCDB zip) is recorded at download time; the corpus build script recomputes them before running; and the sealed-case list (below) is itself committed as a SHA-256 of its JSON serialization. The intent is that a skeptical reader can verify, without trusting the author, that the data evaluated in Section 5 is the data described here—and that nothing was edited between the freeze and the test.

Coverage statistics are reported with their failures rather than their successes: 22 cases carry no SCDB match at all (mostly original-jurisdiction and consolidated matters under variant docketing), 27 additional cases match but have no coded issue area, and audio coverage, while high (98.6%), is absent for a handful of late-window matters. These gaps flow through the pipeline as explicit missing-value categories rather than silent deletions, and the vote analysis of Section 5 simply operates on the rows where a coded direction exists (4,569 of an expected 4,730 vote slots).

3.3 The bench

Thirteen justices appear on the bench in the window, spanning three natural cohorts: the five-justice block that opened the window (Roberts, Scalia, Kennedy, Thomas, Ginsburg), the mid-window joiners (Sotomayor, Kagan, Gorsuch, Kavanaugh, Barrett), and the closers (Breyer’s long service ending in 2021, Jackson arriving in 2022). The 2020 death of Ginsburg and the two 2020 arrivals mean that no single nine-member configuration spans the whole window; vote-level analyses therefore condition on the justices actually sitting, and pairwise agreement (Section 5) is computed on common cases only, keeping the 60 pairs with at least 50 shared votes. Scalia sits for only ten coded votes (he died in February

2016) and is excluded from the pairwise matrix by the $n \geq 50$ rule—retained everywhere else with his small-sample intervals displayed honestly.

4 Protocol: pre-registration and the sealed holdout

4.1 Split and baselines

The experiment fixes its evaluation before fitting anything. The corpus splits temporally: OT2015–OT2019 (2,314 coded votes) is the training window, and OT2020–OT2023 is the evaluation window. All hyper-parameter selection for the challengers of Section 5 uses grouped 5-fold cross-validation *inside the training window only*, with the group being the case — votes of the same case never straddle a fold boundary, because votes within a case are correlated by construction and leaking them across folds silently inflates validation accuracy.

On this skeleton, five statistical baselines were computed and committed before the first challenger was trained:

- **B1—majority class**: predict the modal decision direction of the training window (liberal, at 51.6% of coded training votes) for every vote.
- **B2—always conservative** (the stronger of the two constant policies on this window).
- **B3—always reverse**: predict that the Court upends the court below; evaluated on the subset with a clean affirm/reverse disposition.
- **B4c/B4—per-justice ideology**: predict each justice’s vote as that justice’s modal direction in the training window, at the case level (majority of simulated votes) and the vote level respectively.
- **B5—pairwise agreement**, a descriptive baseline that calibrates how differently the justices actually vote.

B4 is the bar: 63.7% of votes on the full evaluation window (Wilson 95% CI [61.3; 65.9]). It deserves emphasis that this bar is not a floor to be cleared trivially—it is the accuracy of a model that knows one bit per justice (their revealed ideological lean) and nothing about the case. Any claim that a system “reads the law” must beat an adversary that cannot read at all.

4.2 Sealing fifty of the seventy-nine

Seventy-nine decisions in the corpus came down 5–4—the narrowest possible majority and, folklore holds, the least predictable ones. Fifty of them are sealed. The selection is deterministic and was committed before any model training: the sorted list of 79 five-four identifiers seeds a Mersenne Twister (seed derived from the SHA-256 of the sorted list), which draws 50; the resulting list is frozen by the SHA-256 of its JSON serialization, published in the repository’s statistics file. Sealed cases are excluded from the training window, from every cross-validation fold, and from the transparent evaluation reported in this paper—they will be touched exactly once, by the final conditions, in the Final Test.

The sealing rule exists to make the headline question unfakeable. The persona condition (below) claims to extract a justice’s decision-relevant voice from their past opinions; such a claim is only meaningful against cases the extraction pipeline never saw, and only meaningful once. A sealed, hash-frozen, pre-registered holdout is the cheapest device that makes retro-fitting detectable: any post-hoc adjustment to the models after the seal is broken leaves a git trail.

4.3 The decisive test

The Final Test compares conditions A–D on the sealed fifty. The pre-registered claim structure is deliberately symmetric: *if* $B > A$ on never-seen future cases, the persona is extractible from public text—judicial

Table 2: The four final conditions, evaluated once on the sealed holdout.

Condition	Description
A—Zero-shot	An open-weight 8B language model reads the case dossier (briefs, oral-argument transcript, question presented) and answers with a vote direction. No learning from this corpus.
B—Persona	The same base model, QLoRA-fine-tuned on one justice’s <i>past</i> opinions, tested on that justice’s <i>future</i> and sealed cases.
C—Context	The same base model with retrieval of similar earlier opinions (RAG), no fine-tuning.
D—Statistical	The pre-registered baselines of Section 4; B4 is the bar.

Table 3: Pre-registered baselines on the evaluation window (OT2020–OT2023). Wilson 95% intervals.

Baseline	Accuracy (%)	95% CI
B1—majority class (liberal)	43.6	[37.2; 50.1]
B2—always conservative	56.4	[49.9; 62.8]
B3—always reverse	60.1	[51.8; 67.9]
B4c—per-justice ideology, case level	55.8	[49.3; 62.2]
B4—per-justice ideology, vote level	63.7	[61.3; 65.9]

subjectivity leaves a measurable footprint. If $B = A$, the justice’s personality is not in the public data, and whatever is predictable is already priced by ideology. Both outcomes are results; the experiment is built so that neither can be quietly discarded.

One more protocol commitment, unusual for machine-learning papers and central to this one: the public site that displays the experiment’s state is forbidden from showing any model output that does not exist. Until a condition is trained, its pages show the statistical baseline’s call, labeled as such. This is not an aesthetic preference; it is a control on the author’s own incentive to narrate success before the seal is broken.

5 Results

5.1 Baselines: the price of admission

Table 3 and Figure 2 report the pre-registered baselines on the evaluation window. The ordering is informative. The majority-class and constant policies (B1, B2) land in the 44–56% band—barely better than their own training-window priors, a first sign that the evaluation window is not an easy era. The reversal heuristic (B3) reaches 60.1% but on the affirm/reverse subset only, with an interval that warns against taking the point estimate seriously. The per-justice ideology rules dominate: 55.8% at the case level (B4c) and **63.7%** at the vote level (B4) on the full evaluation window. On the transparent (sealed-excluded) subset used for challenger comparison below, B4 scores 63.1% ($n = 1,522$, CI [60.6; 65.5]).

5.2 Pairwise agreement

The descriptive baseline B5 measures how often two justices vote the same way on their common cases. The extremes frame the scale: Kavanaugh and Roberts agree on 94.6% of their 333 common coded votes (CI [91.6; 96.8]), while Sotomayor and Thomas agree on only 54.3% of 518 (CI [50.0; 58.5])—statistically indistinguishable from two justices flipping independent coins that share a slight bias. Figure 3 shows the full matrix, ordered by mean agreement with the conservative anchor block (Thomas and Alito). The block structure is visible without statistics: the modern conservative five form a dark square in one corner, the liberal three-plus in the other, and the two red-outlined extremes are, on inspection, the two ends of the Court’s ideological span. This matrix is the empirical picture of what B4 compresses into one bit per justice.




Figure 2: The number to beat. Wilson 95% intervals for the five baselines on the test window; B4 (red) is the vote-level bar of 63.7%.

Table 4: M3a challengers vs. the bar, transparent test (sealed-excluded; $n = 1,522$ rows common to all rows where B4 is defined). McNemar exact test against B4’s paired predictions.

Model	CV (%)	Test (%)	95% CI	p
M3a-LR (additive)	56.2	58.6	[56.1; 61.1]	0.002
M3a-GB (boosting)	56.1	59.8	[57.3; 62.2]	0.009
M3a-IX (interactions)	57.1	60.4	[57.9; 62.8]	0.047
B4 (bar)	—	63.1	[60.6; 65.5]	—

5.3 M3a: the structured challengers—and the bar holds

The first trained condition, M3a, asks whether the structured, pre-decision attributes of a case—issue area, lower-court disposition, term, and oral-argument duration, joined with justice identity—predict vote direction better than the one-bit ideological rule. Three challengers were trained on the sealed-excluded training window with grouped cross-validation for all hyper-parameter choice: an additive logistic regression (M3a-LR), the same with explicit justice $\times$ issue interactions (M3a-IX), and a histogram gradient-boosting classifier that learns interactions natively (M3a-GB).

Table 4 and Figure 4 give the result: every challenger lands between 58.6% and 60.4% on the common rows, and every comparison against B4 favors the bar (McNemar exact $p \leq 0.047$). The failure is not an artifact of regularization strength or model family; it is robust across three families and survives the interaction model that the site’s per-justice domain-divergence displays would suggest. It is also not a power issue in the usual direction—the challengers are significantly *worse*, not merely insignificantly different.

The ablation in Table 5 locates the failure precisely. Removing justice identity from the feature set collapses cross-validated accuracy from 56.2% to 47.6%—below the class base rate—while removing any other feature group *improves* it (issue area is the most harmful to keep, at +1.7 points when dropped). In other words, on this window, *who* votes is signal; *what the case is about* is noise. The structured challenger did not fail because the metadata was too thin; it failed because the only stable generalizable structure in the data is the attitudinal bit that B4 already uses.

The case-level reading is the same. Simulating each case’s outcome as the majority of the challenger’s predicted votes yields 58.3% (CI [51.5; 64.9]) on the transparent subset—above the case-level ideological baseline (55.8%), but the interval is wide enough that the honest summary is “not distinguishable”, and the vote-level failure stands regardless.




Figure 3: Vote-level agreement across the 60 pairs with at least 50 common coded votes (values are percent agreement). Darker cells are higher agreement; the diagonal is undefined. Red outlines mark the closest (Kavanaugh–Roberts, 94.6%) and widest (Sotomayor–Thomas, 54.3%) pairs. Ordering is by mean agreement with the conservative anchor block.

5.4 What the null prices

Read together, the results give the experiment its current shape. The predictable component of a justice’s vote on this window is, to the resolution that nine terms of structured data can detect, the justice’s ideological direction—exactly what the attitudinal model predicts and exactly what B4 encodes. We phrase the implication precisely: the *tested* structured features add no stable signal beyond the ideological bit on this window. This does not price *all* non-textual information—untested structured sources (parties, counsel, finer temporal or institutional variables) or richer encodings of the same file could in principle hold residual signal; what our result establishes is that the standard SCDB-style coding of the case file does not. The cheapest rich representation that remains untested is the *text*—the briefs, the transcript, the opinion language itself. That is where the pre-registered conditions A–C operate, and the M3a null result is what makes them the critical path rather than a side quest: the structured ceiling is now empirically priced for this representation, so any future claim of textual signal is a claim against a known, pre-registered number.

6 Discussion

6.1 Who votes beats what the case is about

The M3a null result is a negative finding with a positive reading. The failure of three model families to convert issue areas, lower-court dispositions, terms, and argument lengths into votes tells us that the




Figure 4: The first challenger fails against the bar. Wilson 95% intervals on the common rows; p -values are exact McNemar tests against B4’s paired predictions.

Table 5: Feature-group ablation (grouped 5-fold CV on the training window, additive logistic regression).

Feature set	CV accuracy (%)
Full (justice, issue, lc, term, duration)	56.2
drop justice identity	47.6
drop issue area	57.9
drop lower-court disposition	56.7
drop term	56.5
drop argument duration	56.1

mapping from case facts to outcomes is either unstable across time or not expressible in these coarse codes. The ablation sharpens this: keeping issue area actively hurt generalization, which is the signature of a signal that shifted between the training and evaluation windows—the mix and meaning of “economic activity” cases in OT2015 is not the mix of OT2022. Justice identity, by contrast, travels perfectly across the split, because it is a property of the voter, not of the era.

This is a modern, pre-registered echo of the attitudinal model’s central claim [3], obtained with none of its theoretical apparatus: one revealed-ideology bit per justice carries more predictive weight than the entire structured case file. It also explains, mechanically, why the earlier literature’s headlines sit above our bar. Katz et al. [2] report 71.9% at the vote level—but over 1816–2013, an era range containing long stretches of 80–90% unanimity that no modern model enjoys; on the modern, 5–4-heavy window, the honest point of comparison is closer to our setting, and the 63.7% bar is what an ideology-only adversary achieves when unanimity is gone.

6.2 What the null makes the text conditions more interesting, not less

Had a structured challenger cleared the bar, the persona experiment would have been partially redundant: some of the textual signal would already be recoverable from codes. The null result does the opposite—it makes the text conditions the natural next probe. But we state the inference carefully: the M3a failure shows that the *tested* structured representation (issue area, lower-court disposition, term, argument duration, and their pairwise interactions) does not beat the ideological rule on this window. It does not show that all non-textual information is exhausted: party characteristics, counsel characteristics, finer-grained temporal features, institutional signals, or richer encodings of the case file could still carry residual signal that our codes and model families did not capture. What the null does establish is narrower and firmer: any future claim of improvement must now be measured against a known, pre-registered adversary, and the cheapest rich representation of the case that remains untested is its text. The persona condition (B) tests the strongest version of that probe — that any residual textual signal is *individual* and *prospective* — and the stylometry lineage [5, 6] gives it a principled footing: style is known to be persistent and hard

to fake, so a fine-tuned adapter is a reasonable instrument for measuring whether style and substance are entangled.

6.3 Calibrating expectations for the Final Test

The sealed fifty are 5–4 decisions, which the folklore of court watching regards as the least predictable class; and the 63.7% bar was computed on the full evaluation window, not on 5–4s specifically. The protocol therefore re-fits the ideological baseline on the sealed-excluded training window and evaluates it on the sealed population itself: the Final Test comparator is *B4 on the 5–4s*, not the general-window figure, and the report will carry both bars side by side (the pre-registered analysis addendum fixes this reading). Because the sealed population is deliberately harder than the window the bar was computed on, the two numbers may differ substantially, and comparing final conditions to the general-window bar directly would overstate their performance. Both facts cut against the persona condition clearing the bar on the holdout, and the honest prior is that the decisive test may return “no difference”. We stress that this is a designed outcome, not a failure state: a tight null on a sealed, pre-registered holdout with per-justice adapters and audited anti-memorization checks would be, to our knowledge, the cleanest public evidence yet that judicial subjectivity, beyond ideology, is not extractible from the written record—a claim that would stand regardless of how much compute a future replicant throws at it.

7 Limitations

Ground truth is a single coding. Our outcome variable—vote direction, liberal vs. conservative — is SCDB’s editorial coding, produced and maintained by a small team applying published conventions [10]. That coding is the field’s standard, but it is one team’s reading of the cases: directional coding of some opinions (especially mixed or narrow ones) is contestable, and the database’s own documentation flags categories with known unreliability. We did not recode anything; all weaknesses of the coding are inherited as measurement error that attenuates every accuracy in this paper, the bar included.

A short, modern window. Nine terms is a small epoch by the standards of the forecasting literature, and it is deliberately the hardest modern stretch: benches changed twice, 5–4 rates are high, and the ideological middle (Kennedy, then Roberts) moved. Consequences are concrete — the training window contains no votes at all for Barrett and Jackson, so challengers receive no signal for them, and the B4 comparison excludes them by the $n \geq 20$ rule. The Final Test report therefore separates *known-justice forecasting* (justices with training history, where B4 and the persona are both defined) from *new-justice forecasting* (Barrett, Jackson: B4 is undefined by construction, the persona is impossible, and coverage falls to the zero-shot and retrieval conditions). These are different questions—a model that cannot know a new justice is not the same failure as one that knows a justice and still misses—and aggregating across that boundary without labeling it would blur both. We consider the window a feature (the numbers describe *this* Court, not a historical average) but it bounds external validity: nothing here transfers automatically to other eras or other courts.

The transparent test is exploratory. The M3a comparison in Section 5 runs on the non-sealed evaluation window. Although every hyper-parameter was selected by grouped cross-validation inside the training window and the three model families were fixed in advance, the transparent test was run once, on data we had already seen while computing baselines; the multiple comparisons across three challengers are not corrected. The sealed fifty exist precisely to absorb this residual degrees-of-freedom problem at the Final Test.

The text collection is complete but untested. The opinion-text corpus (M1.5) is now complete: all 1,778 targeted opinion texts were collected, normalized, and deduplicated to 793 distinct texts (the remainder are documented textual duplicates across docket aliases), with a zero-leak audit of fourteen

checks committed and re-executed on every push. Every textual condition nevertheless remains untested pending the free-tier GPU training of the persona adapters, and this paper’s null result bounds only the *structured* world. We regard the text as the live question, and we flag the risk we cannot exclude: the eventual persona adapters may overfit memorized opinion content rather than transferable style — mitigations (min- k % probes, cloze audits, temporal splits per justice) are pre-committed in the repository protocol.

One rater of the pipeline. The corpus fusion, the docket matching, and the sealed selection were implemented and audited by a single independent researcher with scripts, hashes, and logs—not by an institutional team. The machinery is public and every claim regenerates from the repository, but the human audit trail is one person deep, and we count that as a limitation of confidence rather than of method.

8 Ethics

The experiment touches only public records: published opinions, docket metadata, and audio of public oral arguments, all released by the Court and mirrored by public-interest archives. No private data, no sealed-at-source material, and no personal data beyond the public professional conduct of public officials is used. The justices are public figures acting in an official capacity; the measured quantities (vote directions, agreement rates, writing style of signed opinions) are matters of public record and public commentary.

Two commitments go beyond data source. First, counterfactual labels are labeled as counterfactuals: when the site or the models display “what the baseline would have voted” for a case a justice never heard (bench turnover makes this common), the row is marked as a simulated call, never presented as a prediction the justice could have made. Second, the project is explicitly non-commercial: no paywall, no product, no data licensing — the corpus, the scripts, the trained adapters, and this paper are published for replication and criticism. The zero-budget constraint is partly ethical in origin: results that can only be reproduced with institutional compute are claims that can only be audited by institutions, and judicial-prediction claims should be auditable by anyone.

A final note on downstream use. A high-accuracy vote predictor would be commercially interesting to lobbyists and litigators; the current state of the project—a 63.7% ideological bar that structured models fail to beat—is, we note dryly, a result with no immediate commercial application, and we are in no hurry to change that.

9 Reproducibility

Every number in this paper regenerates from the public repository <https://github.com/Vitalchiffe/legally-subjective> with three commands and no credentials. The corpus rebuild chain is `scripts/build_corpus.py` (raw sources → frozen corpus with recomputed hashes); the baselines are `scripts/m2_baselines.py`; the M3a challengers, including the grouped cross-validation, the McNemar tests, and the ablation, are `scripts/m3a_train.py`. The paper figures are produced by `scripts/paper_figures.py`. All randomness is seeded and the seeds are printed in the outputs; re-running the M3a script reproduces the numbers of Table 4 exactly.

Provenance is structural, not narrative. The SHA-256 of every raw source file is recorded at download; the frozen corpus recomputes and checks them; the sealed fifty are identified by a deterministic seeded draw whose seed derivation and resulting list hash are both committed. A reader who wishes to verify that the sealed list was not regenerated after training can do so from the git history alone. The dataset rule, the known exclusions (ghost docket groups, consolidated matters without SCDB matches), and the coding caveats are documented in the repository’s docs/ rather than hidden in prose.

The opinion-text corpus (M1.5) is complete: 1,778 collected texts, normalized and deduplicated to 793 distinct documents with full provenance, committed with a fourteen-check zero-leak audit that the repository’s continuous integration re-executes on every push. The remaining artifact under construction is the persona training itself (M3b), running on free-tier GPU sessions; its completion triggers the pre-committed text conditions—at which point this paper’s structured null becomes the fixed background against which the persona question is finally asked.

10 Conclusion

This working paper filed the first half of a pre-registered experiment on the predictability of Supreme Court votes. The frozen corpus exists, is hash-chained, and is public; the statistical baselines are committed, and the strongest of them — one bit of revealed ideology per justice—sets a vote-level bar of 63.7% on the held-out window; and a trained challenger, armed with everything the structured case file knows before the decision, failed to clear the bar with three model families and a clean ablation explaining why. The predictable part of a modern justice’s vote, to the resolution that the tested metadata can see, is the justice.

The second half is where the project’s real question lives. The same corpus—its 793 distinct opinion texts collected, normalized, and audited— will face the sealed fifty with a zero-shot reader, a retrieval reader, and nine per-justice persona adapters trained on each justice’s own past words—each condition pre-registered, each adapter audited for memorization, and each result, positive or null, filed as a finding. The experiment was designed so that either answer will be worth reading: a persona that beats the zero-shot reader on unseen future cases would be, to our knowledge, the first sealed-holdout evidence that a justice-specific textual signal carries vote information; a persona that does not would be, to our knowledge, the cleanest evidence yet that, beyond ideology, the written record keeps its counsel. The bar is set, the seal is frozen, and the remaining variable is the text.

References

- [1] Theodore W. Ruger, Pauline T. Kim, Andrew D. Martin, and Kevin M. Quinn. The supreme court forecasting project: Legal and political science approaches to predicting supreme court decision-making. *Columbia Law Review*, 104(4):1150–1210, 2004.
- [2] Daniel Martin Katz, II Bommarito, Michael J., and Josh Blackman. A general approach for predicting the behavior of the supreme court of the united states. *PLOS ONE*, 12(4):e0174698, 2017.
- [3] Jeffrey A. Segal and Harold J. Spaeth. *The Supreme Court and the Attitudinal Model Revisited*. Cambridge University Press, Cambridge, 2002.
- [4] Andrew D. Martin and Kevin M. Quinn. Dynamic ideal point estimation via Markov chain Monte Carlo for the U.S. supreme court, 1953–1999. *Political Analysis*, 10(2):134–153, 2002.
- [5] Frederick Mosteller and David L. Wallace. *Inference and Disputed Authorship: The Federalist*. Addison-Wesley, Reading, MA, 1964.
- [6] Patrick Juola. Authorship attribution. *Foundations and Trends in Information Retrieval*, 1(3): 233–334, 2006.
- [7] Tim Dettmers, Artidoro Pagnoni, Ari Holtzman, and Luke Zettlemoyer. QLoRA: Efficient fine-tuning of quantized LLMs. In *Advances in Neural Information Processing Systems*, volume 36, 2023.
- [8] Aaron Grattafiori et al. The llama 3 herd of models. *arXiv preprint arXiv:2407.21783*, 2024.

- [9] Free Law Project. CourtListener: U.S. case law, dockets, and oral-argument archive. <https://www.courtlistener.com>. Bulk files of 2026-06-30; accessed 2026-08-28.
- [10] Harold J. Spaeth, Lee Epstein, Andrew D. Martin, Jeffrey A. Segal, Theodore J. Ruger, and Sara C. Benesh. *Supreme Court Database, version 2025 Release 01*. Washington University in St. Louis, St. Louis, MO, 2025. URL <http://scdb.wustl.edu>.