

Abstract

Title	Legally Subjective: How Much Vote Predictability Survives a Strong Ideological Baseline?
Authors	Amine Harch El Korane · Ghali El Alj
Affiliation	Lycée français international Louis-Massignon, Bouskoura–Casablanca, Morocco
Manuscript	Working paper, revision 3 · Corpus-Monde v1 (frozen 2026-08-28) · declarations added 2026-10-09
Replication	github.com/Vitalcheffe/legally-subjective (release tags a1-freeze, m4-freeze, m4-exam)
Length	23 pages · 13,170 words (text + footnotes + references)

Abstract

We ask how predictable U.S. Supreme Court decisions are from public information, and whether modeling individual justices improves that predictability. We assemble a frozen, hash-chained corpus of 569 argued cases (OT2015–OT2023, 13 justices, 4,569 coded votes) and fix the price of admission with pre-registered baselines: a per-justice ideological rule that reaches 63.7% of votes on the held-out window, a bar that three families of structured-feature challengers fail to clear. The pre-registered test of the written record is a sealed exam on fifty hash-frozen 5–4 decisions—zero-shot, persona (biographical prompt plus a QLoRA adapter fine-tuned on each justice's past opinions), retrieval, and the re-fit ideological prior—with predictions frozen before any answer key existed and scored exactly once. Personalization does not detectably improve prediction of the Court's hardest cases: on the primary endpoint, case-direction accuracy, the persona scores 59.18% against the zero-shot's 61.22% (exact McNemar, $p=1.00$). Even the ideological prior collapses to a coin flip at case level on these one-vote-margin decisions (50.0%), while at vote level it still dominates every language condition (75.9% against 50–60%). Personalization does, however, change the model: it moves individual votes (nominally, $p=.047$, unadjusted) without improving case-direction prediction; every simulated justice votes more conservatively than the justice it imitates; the biographical prompt, not the fine-tuned adapter, does most of the moving; a pre-registered stylistic account of persona quality is falsified ($\rho=-0.72$); and personas misvote on the very training cases whose texts they absorbed. Taken together, these results provide no evidence that modeling how an individual justice writes adds detectable predictive power for close case outcomes beyond the information captured by the ideological baseline. All code, frozen artifacts, and sealed predictions are public.

Keywords

judicial prediction; Supreme Court; pre-registration; negative result; large language models; personalization; attitudinal model; open data.

Suggested citation of the working paper. Amine Harch El Korane & Ghali El Alj, *Legally Subjective: How Much Vote Predictability Survives a Strong Ideological Baseline?* (working paper, revision 3, Oct. 2026), github.com/Vitalcheffe/legally-subjective.