

AI-Assistance Disclosure

This is the standalone version of the disclosure printed in the manuscript (Declarations, revision 3). It is written to satisfy the strictest reading of the venue policies concerned: where a policy distinguishes between AI used for background research or copy-editing (often exempt from declaration) and AI used to produce text appearing in the submission (declarable), this document declares the declarable category explicitly.

Statement (verbatim from the manuscript)

This research was conducted with the help of large-language-model coding and writing assistants, used as instruments under the authors' direction and supervision. The authors specified the experimental design decisions, the acceptance criteria, and the analyses to run; the assistants produced draft code, draft text, and documentation that the authors reviewed, tested, corrected, and take full responsibility for. No reported number was accepted on an assistant's word: every figure regenerates from the committed scripts and frozen artifacts described in the reproducibility section of the manuscript, and the sealed-exam chain ran under the authors' control with a lock that is independently verifiable from the public repository.

What this means, category by category

Category of use	Disclosed here
Research assistance (locating and summarizing background materials)	Yes — declared
Code-related assistance (draft code, debugging, documentation)	Yes — declared
Structural analysis and outlining	Yes — declared
Drafting assistance (draft text that appears, in revised form, in the manuscript)	Yes — declared
Copy-editing assistance	Yes — declared
Factual, statistical, methodological, and citation content	Independently verified by the authors against primary sources (court opinions, docket records, the project's own frozen artifacts)

The disclosure deliberately covers the declarable categories; nothing is left implicit.

Verification, not trust

The control on AI assistance is structural rather than narrative. Every number reported in the manuscript regenerates from committed scripts over frozen artifacts, with seeds printed in the outputs; the sealed examination ran under a hash lock (results/m4_exam.lock) that records both the repository head and the SHA-256 of the frozen predictions; and the repository's continuous integration re-executes the fourteen-check zero-leak audit on every push. A reader therefore does not need to trust either the authors or their tools: the chain from corpus to conclusion is checkable, command by command, from the public repository.

Signature. The authors take full responsibility for the entire content of the submission.

Amine Harch El Korane · Ghali El Alj — [date at submission]