

Short Abstract (JOLT)

Manuscript	Legally Subjective: How Much Vote Predictability Survives a Strong Ideological Baseline?
Type	Article submission — Harvard Journal of Law & Technology (via Scholastica)
Authors	Amine Harch El Korane · Ghali El Alj
Length	23 pages · $\approx$ 12,800 words body + footnotes (13,170 with references)

When a large language model is personalized to imitate an individual decision-maker, does the persona add predictive information? We report a pre-registered, sealed test on fifty hash-frozen 5–4 U.S. Supreme Court decisions, drawn from a frozen corpus of 569 argued cases and 4,569 coded votes (OT2015–OT2023). Four conditions predicted the sealed cases blind, with predictions hash-locked before any answer key existed: zero-shot prompting; a persona (biographical prompt plus a QLoRA adapter fine-tuned on each justice’s past opinions); retrieval; and a re-fit ideological prior. Personalization showed no detectable improvement — the persona scores 59.18% case-direction accuracy against the zero-shot’s 61.22% (exact McNemar $p = 1.00$) — while the ideological prior still predicts 75.9% of individual votes yet collapses to chance at case level on these one-vote-margin decisions. The Article analyzes what personalization changes in a model without changing what it predicts — individual votes move, every simulated justice shifts conservative, the biographical prompt rather than the adapter does most of the moving — and what that implies for AI-assisted legal analysis. The full replication package (corpus, scripts, adapters, frozen artifacts, sealed predictions) is public.