

Letter of Introduction

To	The Editors, Harvard Journal of Law & Technology
Re	Article submission — “Legally Subjective: How Much Vote Predictability Survives a Strong Ideological Baseline?”
From	Amine Harch El Korane · Ghali El Alj
Date	[DATE AT SUBMISSION]

Dear Editors,

We are pleased to submit “Legally Subjective: How Much Vote Predictability Survives a Strong Ideological Baseline?” for consideration as an Article in the Journal. The manuscript confronts a question squarely in the Journal’s territory — the intersection of artificial intelligence, big data, and the law: when a large language model is personalized to imitate an individual legal decision-maker, does the personalization actually add predictive information?

We fine-tuned QLoRA persona adapters on the past opinions of thirteen U.S. Supreme Court justices and tested them, in a pre-registered sealed examination on fifty hash-frozen 5–4 decisions, against a zero-shot prompt, a retrieval condition, and a re-fit ideological prior — with all predictions frozen and hash-locked before any answer key existed and scored exactly once. On the primary endpoint, the persona condition scores 59.18% case-direction accuracy against the zero-shot’s 61.22% (exact McNemar, $p = 1.00$); the ideological baseline still predicts 75.9% of individual votes on these cases while collapsing to a coin flip at case level. The Article documents what the persona *does* change — individual votes move, every simulated justice votes more conservatively than the justice it imitates, and the biographical prompt rather than the fine-tuned adapter does most of the moving — and what it does not: any detectable improvement in predicting outcomes.

Because the Journal asks authors to enable verification of their analyses, we direct the editors to the complete replication package: corpus, scripts, persona adapters, frozen artifacts, sealed predictions, and audit reports are public at github.com/Vitalcheffe/legally-subjective, with release tags a1-freeze, m4-freeze, and m4-exam; every CPU-side number regenerates with three commands and no credentials. A replication README accompanies this submission and describes the five-minute verification path. The manuscript runs approximately 12,800 words of body and footnotes (13,170 including references) — below the typical 20,000–30,000-word range of JOLT Articles; we are glad to expand the empirical and doctrinal discussion should the editors prefer.

Declarations accompany the submission: joint design with divided leadership (computational and legal); an AI-assistance disclosure written to be compatible with the Journal’s policy; no funding; no competing interests; and full data and code availability. This research was conducted independently by the two authors, who are students at the Lycée français international Louis-Massignon (Bouskoura–Casablanca, Morocco); the school is named for identification only and does not endorse the research.

We would be honored by the Journal’s consideration, and we thank the editors for their time.

Respectfully submitted,

Amine Harch El Korane · Ghali El Alj

[EMAIL À COMPLÉTER] · [EMAIL À COMPLÉTER]

SOURCE: Enclosures: manuscript (Word), short abstract, curriculum vitae of each author, AI-assistance disclosure, replication README.