

# CACM — Pitch

|                       |                                                                                           |
|-----------------------|-------------------------------------------------------------------------------------------|
| Proposed contribution | Contributed article / viewpoint, computing-methods angle                                  |
| Working title         | When Personalization Fails: Testing Whether LLM Personas Actually Predict Human Decisions |
| Authors               | Amine Harch El Korane · Ghali El Alj                                                      |
| Contact               | [EMAIL À COMPLÉTER]                                                                       |
| Length target         | ≈ 1,500–3,000 words (to be set with the editors)                                          |

## The pitch, in one paragraph

*Personalizing an LLM on a decision-maker’s own writings is sold as prediction, but it is almost never tested as prediction. We sealed the test: a frozen, hash-chained corpus of 569 U.S. Supreme Court cases and 4,569 coded votes; a pre-registered ideological baseline that already predicts 63.7% of votes; and a sealed exam on fifty 5–4 decisions predicted blind by a zero-shot prompt, a QLoRA persona (biographical prompt plus an adapter fine-tuned on each justice’s past opinions), retrieval, and the re-fit prior — with predictions hash-locked before the answer key existed. The persona showed no detectable improvement (59.18% vs 61.22% case-direction; exact McNemar  $p = 1.00$ ) while systematically shifting every simulated justice conservative and compressing the ideological spectrum into a narrow band. A pre-registered decomposition shows the biographical prompt, not the fine-tuned adapter, does most of the moving. Everything — corpus, scripts, adapters, frozen artifacts, sealed predictions — is public and regenerates with three commands.*

## Why CACM readers care

- Evaluation methodology, not law: baseline-first claims, sealed tests, leakage audits, and decomposition of personalization into prompt vs weights — reusable wherever LLMs are personalized to imitate humans.
- A clean, reproducible negative result with a mechanism: compression and conservative drift, quantified seat by seat.
- Reproducibility as infrastructure: hash-locked predictions, CI-replayed leakage audits, three-command regeneration — a template readers can adopt for their own evaluations.

## Why now, why us

Persona-style personalization is being deployed in consumer and enterprise products faster than it is being evaluated; a sealed, pre-registered null from a fully public pipeline is the counterweight the conversation lacks. We built the entire experiment ourselves — corpus, adapters, seal, audits — at zero cost on public data and free-tier compute, which is itself part of the story: rigorous evaluation of LLM claims no longer requires institutional resources.

## What we can provide

- The computing-oriented draft (CACM\_ARTICLE.pdf,  $\approx$  3,200 words, four figures ready to reuse).
- The parent manuscript (23 pages) as background.
- A live five-minute verification path through the public repository ([github.com/Vitalcheffe/legally-subjective](https://github.com/Vitalcheffe/legally-subjective)).
- Interviews and demo material on request.

**Route note (source discipline).** The exact editorial route for unsolicited contributions — and the current word-count windows for contributed articles vs viewpoints — is to be confirmed on CACM's official contributor pages at pitch time; anything not confirmed there is treated as NOT ESTABLISHED BY OFFICIAL SOURCE. Do not attach the 23-page law-review manuscript to the first approach: the pitch is the product; the draft is the proof of depth.

## Companion files

- SUBMISSIONS/07\_CACM/CACM\_ARTICLE.pdf — the computing-oriented draft (Product C).
- SUBMISSIONS/00\_MASTER/MASTER\_PAPER.pdf — parent manuscript (Products A/B).
- SUBMISSIONS/08\_MEDIA/ — press preparation (Product D angle, distinct audience).